

Mapping With Melanin — Integrity, Accountability & Partnership Framework

Product Specification Document

1. Executive Summary

Mapping With Melanin (MWM) is a community-powered platform designed to help minorities find trusted businesses, stay safe, and connect with their community. This document outlines the comprehensive product specification for the platform's Integrity, Accountability, and Partnership Framework. It details the mechanisms for business classification, safety reporting, legal protections, partnership pipelines, diaspora contributions, and AI personalization, ensuring a safe, equitable, and legally sound environment for all users and listed businesses.

2. Business Classification Tiers

The platform categorizes businesses into distinct tiers to ensure minority-owned businesses are prioritized while providing a pathway for community-verified allied businesses.

Tier 1: Minority-Owned Businesses (PRIMARY)

This tier represents the core focus of the MWM platform.

- **Visibility:** Always shown first in search results and recommendations.
- **Features:** Full profile access, community reviews, and Kinfolk AI recommendations.
- **Opportunities:** Eligible for flash deals, creator partnerships, and demand intelligence insights.
- **Identification:** Ownership badge prominently displayed.
- **Engagement:** Actively invited and supported by the platform.

Tier 2: Allied Businesses (PASSIVE)

This tier is a critical distinction, representing non-minority-owned businesses verified as safe by the community.

- **Definition:** Non-minority-owned businesses that the community has verified as safe spaces (e.g., "Hyatt Wisconsin — Community Verified Safe Space").

- **Visibility Constraints:** NEVER promoted above a minority-owned alternative in the same category.
- **Onboarding:** NEVER contacted or alerted when first added to the platform. They remain PASSIVE entries until accumulating sufficient community engagement data.
- **Partnership Pathway:** Only AFTER sufficient tracking data (e.g., 500+ click-throughs from MWM) do they receive a partnership opportunity link. The pitch comes TO THEM based on demonstrated value.
- **Accountability:** Subject to immediate deprioritization if discrimination is reported. Can lose their listing entirely if community trust erodes.
- **Identification:** Distinct "Community Verified" badge, contrasting with the ownership badge.

3. Safety Reporting & Reporter Protection

Protecting users who report concerns is paramount. However, there is a **critical distinction** between two fundamentally different types of reports, each requiring different response protocols, timelines, and protection levels.

Report Type A: Community Safety Concerns (Visitor/Customer Experience)

These reports protect community members as **customers, visitors, or passersby**. The reporter has a transient relationship with the business — they visited once, experienced something, and left.

Examples: "I was followed in the parking lot" / "Staff used a racial slur" / "They refused to serve us" / "I felt unsafe here" / "This gas station has a history of discrimination"

Response Protocol for Type A:

Attribute	Type A — Safety Concern
Reporter relationship to business	Transient (customer/visitor)
Reporter retaliation risk	Low (they don't return daily)
Aggregation threshold	2 independent reports within 60 days
Time delay before action	7 days (shorter — community needs to know quickly)
Visibility to community	Affects public-facing trust score

What community sees	"Community safety concern reported" on the business listing
Purpose	Protect the NEXT person who might visit
Graduated response speed	Faster — safety is urgent

Type A reports directly affect the business's community trust score and can trigger the graduated response system (Levels 1-5) more quickly because the priority is **preventing harm to the next visitor**. The reporter is relatively safe because they have no ongoing relationship with the business.

Report Type B: Unfair Labor & Workplace Practices (Employee Experience)

These reports protect community members as **employees or job applicants**. The reporter has an ongoing, daily relationship with the business — they go back to that workplace every day. This creates an extreme retaliation risk.

Examples: "They pay below minimum wage" / "Hostile work environment for Black employees" / "They fired someone for wearing natural hair" / "Retaliation against employees who speak up" / "Discriminatory promotion practices"

Response Protocol for Type B:

Attribute	Type B — Labor/Workplace Report
Reporter relationship to business	Ongoing (employee/applicant)
Reporter retaliation risk	EXTREME (they return daily, can be fired)
Aggregation threshold	3+ independent reports over 90+ days
Time delay before action	30-60 days minimum (maximum protection)
Visibility to community	Does NOT affect public-facing trust score initially
What community sees	Private warning to members considering employment: "Community members have reported workplace concerns"
Purpose	Warn potential employees WITHOUT endangering current ones
Graduated response speed	Slow and deliberate — reporter safety is paramount

Type B reports are handled with **maximum reporter protection** because the reporter is trapped in a daily relationship with the business. The system must ensure that even if only one person from a small team files a report, they cannot be identified through temporal correlation, team size analysis, or any other inference.

Separation of Scores

Businesses maintain TWO independent trust dimensions:

1. **Visitor Safety Score** — informed by Type A reports. Visible to all community members browsing/searching. Affects search ranking and recommendations.
2. **Employer Trust Score** — informed by Type B reports. Visible ONLY to community members who are actively searching for jobs or who ask Kinfolk AI about working somewhere. Does NOT affect the business's visibility as a place to visit/shop.

This separation is critical: a restaurant might be a wonderful place to eat (high visitor safety) but a terrible place to work (low employer trust). These are independent assessments that serve different community needs.

The Critical Problem

If an employer is consistently ranked negatively for workplace discrimination, the system MUST NOT allow the report to be traced back to a specific employee. Consider: if there is one Black gay woman in a 50-person office, and a workplace discrimination report appears the week she started, she could be identified and retaliated against — even fired.

Protection Mechanisms (Apply to BOTH types, with stricter thresholds for Type B)

1. **Anonymous Aggregation Threshold:** No action is taken on a single report. Type A requires 2 independent reports within 60 days. Type B requires 3+ independent reports over 90+ days. This prevents temporal correlation.
2. **Time-Delayed Visibility:** Reports are collected, but their effects are NOT immediately visible. Type A: 7-day delay. Type B: 30-60 day delay. The longer delay for Type B prevents anyone from correlating "a report appeared right after Person X started."
3. **Geographic Fuzzing:** For Type B (labor) reports, concerns are aggregated at the brand/chain level, not individual location level, until the threshold is high enough. "[Brand] has received workplace concerns" — not "[Brand] at [specific address]." Type A reports can be location-specific once threshold is met because visitor safety is location-dependent.

4. **No Individual Attribution:** The platform NEVER reveals how many reports exist, who reported, when reports were filed, or any metadata that could identify reporters. The only visible output is the aggregated score change.
5. **Reporter Identity Separation:** Reports are stored with a one-way hash of the reporter's ID. Even database administrators cannot easily correlate reports to specific users without a deliberate forensic query.
6. **Retaliation Detection:** If a user's account is reported or flagged shortly after they filed a safety report, the system flags this as potential retaliation for administrative review.
7. **Team Size Protection (Type B only):** If the system detects that a business has fewer than 10 community members associated with it as employees (based on profile data), Type B reports require an even HIGHER threshold (5+ reports) before any action, because small teams make individual identification easier.

Report Submission Flow

When a community member submits a report, the first question is:

"How did you experience this?"

- "I visited or shopped here" → Type A (Safety Concern)
- "I work here or applied for a job here" → Type B (Workplace Report)

This single question determines the entire protection protocol. Type B reporters see an additional reassurance message: "Your report is protected by our maximum-security protocol. It will never be visible until multiple independent community members share similar experiences over an extended period. Your identity is cryptographically separated from your report."

4. Legal Protection Framework

To mitigate the risk of defamation lawsuits from deprioritized businesses, MWM must operate strictly as a platform, not a publisher.

Defamation Mitigation Strategy

1. **Platform vs. Publisher Distinction (Section 230):** MWM is a platform for community-reported experiences, NOT making editorial judgments. Language must reflect this (e.g., "Community members have shared concerns" instead of "This business discriminates"). All visible safety information is attributed to community consensus.
2. **Terms of Service Requirements:** All businesses listed agree to community feedback visibility. Allied businesses are listed as community resources, not commercial listings.

A clear disclaimer must be present: "Trust scores reflect aggregated community experiences and do not represent the views of Mapping With Melanin, Inc."

3. **No Defamatory Statements:** Never use words like "discriminates," "racist," or "unsafe" as platform-generated labels. Use neutral, factual language like "Community trust score: Low" or "Community concerns reported." Allow businesses a right of reply.
4. **Graduated Response System:**
 - **Level 1:** No visible change (collecting data).
 - **Level 2:** Subtle deprioritization in search results (not visible to the business).
 - **Level 3:** "Community concerns noted" badge appears (business notified, can respond).
 - **Level 4:** Business removed from recommendations but still searchable.
 - **Level 5:** Business delisted (only for severe/verified cases, with human review).
5. **Human Review Gate:** No business is publicly flagged or delisted without human (admin) review. Automated systems only handle Levels 1-2. Levels 3-5 require founder/admin approval.
6. **Documentation Trail:** Every trust score change is logged with the aggregated reason, providing a legal defense based on acting on aggregated community feedback following published policy.

5. Partnership Pipeline (for Allied Businesses)

This pipeline defines the journey from a passive community listing to an active partner.

Stage	Name	Description	Criteria
Stage 1	Community Addition (Day 0)	Added by community; business is not contacted.	3+ independent nominations.
Stage 2	Data Accumulation (Days 1-90)	Tracking views, clicks, saves, and mentions.	N/A (Data gathering phase).
Stage 3	Partnership Threshold Reached	System flags for admin review based on engagement.	500+ click-throughs OR 100+ saves OR 50+ direction requests.
Stage 4	Outreach (Admin-initiated)	MWM sends a partnership pitch based on proven value.	Admin approval required.
Stage 5	Active Partner	Access to analytics dashboard, ability to	Paid partnership fee; must maintain

		respond to feedback.	community trust.
--	--	----------------------	------------------

Note: Active partners are still NEVER promoted above minority-owned alternatives.

6. Diaspora International Business Contributions

This feature allows members to add trusted businesses from their home countries or travels.

Contribution Flow

1. **Identification:** Member identifies as part of the diaspora during onboarding.
2. **Prompt:** In-app prompt encourages adding safe spaces in specific countries.
3. **Data Capture:** Form captures business name, location, relationship (owner, family, visitor, community recommended), safety context (e.g., LGBTQIA+, Black travelers), and cultural context.
4. **Trust Weighting:**
 - "I own this" → Highest trust (verified via claim flow).
 - "Family owns this" → High trust.
 - "I've visited" → Medium trust.
 - "Community recommended" → Requires 2+ independent confirmations.
5. **Privacy Considerations:**
 - Safety tags (e.g., LGBTQIA+) may only be visible to users who self-identify similarly in regions where such labels pose risks.
 - Contributors can remain anonymous.
 - The system eventually seeks owner confirmation, but listings can exist based on community trust.

7. Onboarding → Kinfolk AI Pipeline

Data collected during onboarding must immediately feed the Kinfolk AI to provide personalized experiences from day one.

Data Utilization Strategy

1. **Immediate Storage:** Onboarding data (safety preferences, travel interests, identity/diaspora designations) is immediately stored server-side.

2. **Profile Creation:** Creates an initial Kinfolk AI profile based on stated preferences.
3. **Day One Personalization:** Kinfolk uses this profile to personalize recommendations immediately.
4. **Continuous Refinement:** Over time, Kinfolk confirms and refines preferences based on behavioral signals (searches, saves, reviews). Example: "We noticed you save a lot of Caribbean restaurants..."

8. Accountability Without Surveillance

The platform must operate on the principle that it learns from what users CHOOSE to share, never from what it secretly observes.

- **Location Data:** Used only when explicitly granted and for stated purposes.
- **Search History:** Feeds recommendations, but users can view and delete their history.
- **AI Transparency:** Kinfolk AI explains the reasoning behind recommendations.
- **User Control:** Users can correct Kinfolk's assumptions.
- **Value Exchange:** All data collection is permission-based with a clear value proposition (e.g., "Share your travel plans → get safety tips for your route").

9. Implementation Priority

Given the existing codebase, the implementation is phased as follows:

Phase 1 (Pre-tour, website only)

- Add "ambassador" as a valid `previewChoice` in the waitlist.
- Develop the partnership tracking dashboard (admin view of `external-clicks` analytics).

Phase 2 (Post-Apple-approval)

- Implement Allied Business tier distinction in the business model.
- Connect deprioritization logic to safety reports.
- Implement the anonymous aggregation threshold system.
- Ensure onboarding data flows to server-side storage and Kinfolk profile creation.

Phase 3 (Post-tour)

- Automate partnership outreach (threshold-based notifications).
- Implement the diaspora international business contribution flow.

- Develop Kinfolk confirmation prompts.
 - Implement the legal framework (ToS updates, response system for businesses).
-

Developer Guidelines (Replit Integration)

This section details how the new requirements integrate with the existing MWM codebase.

1. Partnership Tracking Dashboard

The existing `/external-clicks` endpoint captures the necessary data (`institutionName` , `institutionType` , `institutionUrl` , `referenceType` , `referenceId` , `source` , `isSafetyRelated` , `city` , `state` , `userId`).

- **Action:** Build an admin dashboard that aggregates this data.
- **Query Logic:** Group by `referenceId` (business ID) where `institutionType` is 'Allied'. Sum the occurrences to track progress toward the Stage 3 thresholds (e.g., 500+ click-throughs).

2. Safety Reporting & Deprioritization

The existing endpoints (`/reports` , `/reports/proximity-warnings` , `/incidents` , `/safety-tips` , `/safety-context`) need to be integrated with the new threshold logic.


- **Action:** Modify the reporting logic to store reporter IDs as a one-way hash.
- **Aggregation Logic:** Implement a background job (or query logic) that checks for the 3+ independent reports within a 30-day window condition before updating a business's aggregated trust score.
- **Visibility Delay:** Add a `visible_after` timestamp to aggregated trust score updates, set to 14-30 days after the threshold is met.

3. Onboarding to Kinfolk AI

Onboarding data (ownership designations, diaspora countries) must be pushed to the server immediately.

- **Action:** Ensure the onboarding flow makes an API call to store this data against the user's profile upon completion, rather than just relying on local `AsyncStorage` .
- **Kinfolk Integration:** The Kinfolk AI service must query this user profile data to initialize its context for recommendations.

4. API Schema Updates (Conceptual)



JSON

```
// Example update to Business Schema to support Tiers and Trust Scores
{
  "businessId": "uuid",
  "name": "string",
  "tier": "Minority-Owned | Allied",
  "trustScore": {
    "level": 1, // 1-5 scale based on Graduated Response
    "lastUpdated": "timestamp",
    "visibleAfter": "timestamp" // Implements Time-Delayed Visibility
  },
  "engagementMetrics": {
    "clicks": "number",
    "saves": "number",
    "directionRequests": "number"
  }
}
```