

Safety Language & Philosophy Audit — Locked Decision

Status: FIX IN BUILD 103 — This is a LOCKED PLATFORM IDENTITY DECISION

The Philosophy (Non-Negotiable)

MWM's safety system is **community-experience-based**, NOT crime-statistics-based. The distinction:

What MWM Safety IS	What MWM Safety IS NOT
How minority community members EXPERIENCE a place	Crime rates from police databases
Do they feel welcome, respected, safe from harm?	"Low crime" = "safe neighborhood"
Community-reported experiences of discrimination, hostility, profiling	Police-reported arrest data that reflects over-policing
Business quality, amenities, job growth, community investment	Property values or income levels as proxies for safety
Incidents of brutality, discrimination, hostility toward minorities	A condemnation of minority communities
"Where does OUR community feel safe?"	"What neighborhoods have low crime?"
Safety FOR us	Safety FROM us

The critical distinction: A wealthy white suburb with a history of racial profiling is LESS safe for our community than a Black neighborhood with strong community ties. Traditional crime data would say the opposite. That's why we don't use it as our primary signal.

CRITICAL PROBLEM FOUND: safety-context.ts

File: artifacts/api-server/src/routes/safety-context.ts

This file pulls LIVE CRIME DATA from police departments:

- Chicago Police (crimes.json)
- NYPD (arrest data)
- Philadelphia Police Department (incidents)
- DC Metropolitan Police (offenses)

This directly contradicts the platform's stated philosophy. The safety-info.tsx page literally says: "We don't infer safety from crime statistics, census data, or third-party databases — all of which carry systemic bias."

Decision: OPTION A — Remove the crime data endpoint entirely.

AUTHORIZED. Delete `safety-context.ts` and all references to it. Safety scores come ONLY from community surveys, member reports, and the neighborhood survey system. The crime data endpoint contradicts the brand promise and is not needed — the community reporting system handles violence reporting.

Remove:

- `artifacts/api-server/src/routes/safety-context.ts` — delete the file
- `artifacts/api-server/src/routes/index.ts` line 11 — remove the import
- `artifacts/api-server/src/routes/index.ts` line 162 — remove `router.use(safetyContextRouter)`
- Any frontend references to `/api/safety-context` endpoint
- The "Live crime data is temporarily unavailable" error message (line 172)

Keep: All community-reported safety features (Safety Tips, Unsafe Space Reports, Anonymous Reports, Neighborhood Surveys). These are the correct data source.

LANGUAGE FIXES NEEDED

Issue 1: "Safe Neighborhoods" Framing

Several places use "safe neighborhoods" which implies some neighborhoods are inherently unsafe (coded language that stigmatizes minority communities).

File	Line	Current	Replace With
<code>cities.tsx</code>	101	"safest neighborhoods"	"most welcoming neighborhoods"
<code>mobile/(tabs)/index.tsx</code>	883	"Spaces for rent, sale & business in safe neighborhoods"	"Spaces for rent, sale & business in

			welcoming communities"
mobile/spaces.tsx	248	"Tell the community about a great space in a safe neighborhood"	"Tell the community about a great space in a welcoming neighborhood"
mobile/onboarding/travel.tsx	87	"safe neighborhoods, cultural events, and community recommendations"	"welcoming neighborhoods, cultural events, and community recommendations"

Issue 2: "Travel While Black/Minority" Framing

The safety page FAQ (line 22) says: "Our safety scores reflect real experiences from people who share your background, not generic ratings." — This is GOOD. Keep it.

But line 12 of safety-info.tsx says: "Safety here means how it actually felt to be there, **as a Black person**, in that space." — This should be updated to be inclusive:

File	Line	Current	Replace With
safety-info.tsx	12	"as a Black person, in that space"	"as a minority, in that space"

Issue 3: Safety FAQ Still Says "Black-Owned"

File	Line	Current	Replace With
safety.tsx	19	"All Black-owned businesses are welcome to join...Any Black-owned business can list..."	"All minority-owned businesses are welcome to join...Any minority-owned business can list..."

(This may have already been caught in the language audit — confirm it was fixed.)

Issue 4: "Melaninated Travelers" vs. "Minority Travelers"

The safety page uses "Melaninated travelers" and "Melanated community" — this is FINE and should stay. It's the platform's identity language. Just ensure it's consistent (sometimes "Melaninated" sometimes "Melanated" — pick one spelling).

Issue 5: "Highest Safety Scores" City List

Line 23 of safety.tsx says: "Our top-rated cities include Atlanta, GA; Houston, TX; Chicago, IL; Miami, FL; and Washington, DC — all with strong Black business communities and high community safety scores."

File	Line	Current	Replace With
safety.tsx	23	"all with strong Black business communities"	"all with strong minority business communities"

WHAT'S ALREADY CORRECT (Do Not Change)

The following are philosophically aligned and should stay:

1. **safety-info.tsx** — "We don't infer safety from crime statistics, census data, or third-party databases — all of which carry systemic bias."  PERFECT
2. **safety-info.tsx** — "When a community member reports a racial profiling incident, harassment, or police involvement...that data is flagged, reviewed, and weighted more heavily in our safety model."  CORRECT — this is community-reported, not police-reported
3. **DemoS06Map.tsx** — "Data from people with lived experience — not crime databases"  PERFECT
4. **DemoS10MapText.tsx** — "Neighborhood color-coding is sourced entirely from diaspora member surveys — not crime statistics, police data, or algorithms that encode historical bias."  PERFECT
5. **DemoS12SafetyHubText.tsx** — "Community Intelligence — not crime data"  PERFECT
6. **DemoS12SafetyHubText.tsx** — "Report Police or ICE activity, track officer transfers, flag unsafe or discriminatory spaces."  CORRECT — reporting police activity is a safety feature FOR the community
7. **Neighborhood Survey** — Uses "welcoming" scale (Very welcoming → Uncomfortable) rather than "safe/dangerous" binary.  CORRECT

8. **welcome.tsx** — "Our neighborhood safety scores are powered by real community reports — not police data."  PERFECT
 9. **safety-hub.tsx** — "Report an Unsafe Space: Flag any business or venue where you experienced unsafe, discriminatory, or unwelcoming treatment."  CORRECT — this is about the community's experience, not crime
-

WHAT SAFETY SCORES SHOULD BE BASED ON (Locked Definition)

MWM Safety Scores are calculated from:

1. **Community Surveys** (primary signal — 60% weight)
 - Daytime welcoming rating
 - Nighttime welcoming rating
 - "Would Return Alone" metric
 - Walkability/transit comfort
 - Atmosphere/cultural inclusion
2. **Incident Reports** (secondary signal — 25% weight)
 - Discrimination experiences
 - Racial profiling
 - Hostile staff/environments
 - Harassment
 - Hate crimes reported BY community members
3. **Business & Community Investment** (tertiary signal — 15% weight)
 - Number of minority-owned businesses in area
 - Community amenities (parks, cultural centers, gathering spaces)
 - Job growth / economic opportunity
 - Access to services (healthcare, education, transit)

NOT included in safety scores:

- Police crime statistics
- Arrest data
- Census income data
- Property values

- Insurance risk assessments
 - Any third-party database that encodes historical bias
-

REGRESSION PREVENTION

Add to `.agents/memory/` a new file `safety-philosophy.md` :

Markdown

```
---
name: Safety Philosophy – Locked Decision
description: MWM safety scores are community-experience-based, NOT crime-statistics
---

## The Rule
Safety on MWM means: "How does our community EXPERIENCE this place?"
Safety on MWM does NOT mean: "What do police crime statistics say about this area?"

## Language Rules
- NEVER use "safe neighborhood" (implies some neighborhoods are inherently unsafe)
- USE "welcoming neighborhood" or "community-trusted area" instead
- NEVER reference crime rates as a measure of safety
- ALWAYS frame safety as community experience, not absence of crime
- Safety scores come from COMMUNITY SURVEYS, not police data

## The Critical Distinction
A wealthy white suburb with racial profiling = LOW safety score for our community
A Black neighborhood with strong community ties = HIGH safety score for our community
Traditional crime data says the opposite. That's why we reject it.
```

NEW FEATURE: Active Emergency Alert System (Build 104 — Pre-Tour)

The Problem

The current Safety Hub has "Submit Safety Tip" which sounds retrospective — like reporting a bad experience from last week. It does NOT clearly communicate that this is where you'd report an ACTIVE SHOOTER, a mass casualty event, or real-time danger. Community members need a way to say "THIS IS HAPPENING RIGHT NOW" and alert everyone nearby immediately.

The Solution

Add a distinct, visually urgent option to the Safety Hub:

" **ACTIVE ALERT — Report an Emergency**"

"Something dangerous is happening RIGHT NOW. Alert community members in this area immediately."

Alert Categories:

Category	Icon	Description
Active violence / shooting		Gunfire, active shooter, mass casualty event
Hate crime in progress		Targeted attack on community members
Civil unrest / dangerous protest		Protest turning violent, tear gas, riot
ICE raid / immigration enforcement		Immigration enforcement activity in area
Natural disaster / severe weather		Tornado, flood, earthquake, severe storm
Infrastructure emergency		Gas leak, bridge collapse, building fire

How It Works:

1. Community member taps " Active Alert" (prominent, red, impossible to miss)
2. They select the category and pin the EXACT location
3. They add a brief description (optional — speed matters more than detail)
4. ALL verified community members within 5 miles get an IMMEDIATE push notification
5. The alert shows on the map as a pulsing red indicator
6. Other members can:
 - Confirm: "I can verify this is happening" (adds weight)
 - Update: "It's moved to [location]" or "Police are on scene"
 - Resolve: "This appears to be over" (after 3+ members confirm resolved, alert fades)

7. Kinfolk proactively messages anyone with bookings, plans, or check-ins in that area:

- "⚠️ There's an active alert near [your planned location]. [Category]. Stay safe — avoid the area if possible."

How It's DIFFERENT from Safety Tips:

Safety Tip	Active Alert
Retrospective (happened in the past)	Real-time (happening NOW)
Informs future decisions	Protects people NOW
No push notification to others	IMMEDIATE push to all nearby members
Appears in safety score over time	Appears on map instantly as urgent indicator
"I had a bad experience here"	"THERE IS DANGER HERE RIGHT NOW"

UI Requirements:

- The Active Alert button should be visually distinct from everything else in the Safety Hub
- Red/urgent color, larger than other options
- Accessible in 2 taps maximum (Safety Hub → Active Alert → pin location → submit)
- The push notification should use iOS Critical Alert style (bypasses Do Not Disturb if possible)
- The map indicator should pulse/animate to draw attention
- Alerts auto-expire after 4 hours unless renewed by community confirmation

Abuse Prevention:

- Only verified members can submit Active Alerts (prevents trolling)
- False alerts flagged by 3+ members → alert removed + submitter warned
- Repeat false alerters lose Active Alert privileges
- Each alert requires location pin (can't alert the whole city without specificity)

Kinfolk Integration:

- If a member has a booking or check-in near an active alert, Kinfolk messages them immediately

- If a member is currently navigating to a location near an alert, Kinfolk interrupts with a warning
 - After an alert resolves, Kinfolk follows up: "The alert near [location] has been resolved. Stay safe out there."
 - Kinfolk can aggregate: "There have been 3 community alerts in [area] this month. Want me to show you alternative routes?"
-

Verification Required

1. Confirm safety-context.ts has been deleted (crime data endpoint removed)
2. Show updated language for "safe neighborhoods" → "welcoming neighborhoods"
3. Confirm safety-info.tsx line 12 updated to "as a minority"
4. Confirm safety FAQ line 19 was caught in the language audit
5. Confirm the safety-philosophy.md memory file was created
6. Confirm: nowhere in the user-facing app does a safety score reference crime statistics as its source
7. Confirm Active Alert feature is in Build 104 scope